CodeInsight: Automated Codebase Health & Technical Debt Prioritization Platform

A report submitted for the course named Project II (CS4203)

Submitted By

KRISHNA ANAND

SEMESTER VII

230101082

Supervised By

DR. RAJKUMARI BIDYALAKSHMI DEVI

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

INDIAN INSTITUTE OF INFORMATION TECHNOLOGY SENAPATI, MANIPUR

NOVEMBER 2026

Declaration

In this submission, I have expressed my idea in my own words, and I have adequately cited and referenced any ideas or words that were taken from another source. I also declare that I adhere to all principles of academic honesty and integrity and that I have not misrepresented or falsified any ideas, data, facts, or sources in this submission. If any violation of the above is made, I understand that the institute may take disciplinary action. Such a violation may also engender disciplinary action from the sources which were not properly cited or permission not taken when needed.

KRISHNA ANAND
230101082

DATE: 5 November 2026

Department of Computer Science and Engineering
Indian Institute of Information Technology Senapati, Manipur

Dr. Rajkumari Bidyalakshmi Devi Email: bidyalakshmi@iiitmanipur.ac.in
Assistant Professor, CSE, IIIT Manipur Contact No: 7005126046

To Whom It May Concern

This is to certify that the project report entitled "**CodeInsight: Automated Codebase Health & Technical Debt Prioritization Platform**," submitted to the Department of Computer Science and Engineering, Indian Institute of Information Technology Senapati, Manipur in partial fulfilment for the award of degree of Bachelor of Technology in Computer Science and Engineering is a record of bonafide work carried out by **KRISHNA ANAND** bearing roll number **230101082** for the course Project II (CS4203).

Signature of the Supervisor

(DR. RAJKUMARI BIDYALAKSHMI DEVI)

Signature of the Examiner 1 .....

Signature of the Examiner 2 .....

Signature of the Examiner 3 .....

Signature of the Examiner 4 .....

Abstract

As software systems grow and undergo rapid iterative enhancements, code quality frequently degrades under continuous delivery pressures, resulting in technical debt. Left unaddressed, technical debt manifests as high cognitive complexity, deeply nested control logic, and widespread copy-paste code duplication, drastically reducing developer velocity and increasing defect rates. Traditional enterprise code inspection suites impose heavy setup friction, requiring complex relational databases, container runtimes, and substantial memory allocations, while lightweight linters provide isolated style warnings without high-level architectural insight.

This project presents **CodeInsight**, an automated codebase health analysis and technical debt prioritization platform engineered for immediate developer utility. CodeInsight pairs a high-performance C++ analysis engine powered by Tree-sitter abstract syntax tree (AST) parsing with a Node.js Express orchestration backend, a zero-dependency in-memory persistence layer, and a high-density React/Vite dashboard. The platform extracts exact function-level metrics—including McCabe cyclomatic complexity, nesting depth, and parameter counts—and identifies multi-file Type-1 code clones through a normalized sliding-window token hashing pipeline with maximal forward block expansion.

To provide clear, actionable guidance, CodeInsight formulates an explainable technical debt risk score:

$$R = 0.45C + 0.35D + 0.20S$$

where structural debt S is decomposed into physical length, nesting depth, and parameter volume ($S = 0.40L + 0.35N + 0.25P$). The platform generates explicit human-readable risk explanations and computes remediation priority tiers ($P1$ to $P4$) to help engineering leads focus refactoring resources effectively. By eliminating external database dependencies, CodeInsight achieves instantaneous startup in under 5 seconds, maintains an operational memory footprint under 120 MB, and achieves an AST parsing and clone detection throughput exceeding 48,000 lines of code per second.

Acknowledgement

I am deeply grateful to my project supervisor, **Dr. Rajkumari Bidyalakshmi Devi**, for her continuous guidance, valuable technical feedback, and encouragement throughout the development of this project. Her insights into software engineering methodologies, formal program analysis, and architectural design were instrumental in shaping the evaluation framework and metric formulations of CodeInsight.

I would also like to express my sincere appreciation to the faculty members and peers at the Department of Computer Science and Engineering, Indian Institute of Information Technology Senapati, Manipur, for their constructive critiques and helpful suggestions during the development and benchmarking phases of this work.

Krishna Anand
230101082

Contents

Declaration **1**

Certificate **2**

Abstract **3**

Acknowledgement **4**

List of Figures **8**

List of Tables **9**

1 Introduction **10**

1.1 Introduction 10
1.2 Motivation 10
1.3 Problem Statement 11
1.4 Objective 12
1.5 What is our Contribution? 12
1.6 Project Timeline 14

2 Literature Survey **15**

2.1 Existing Systems and Methodologies 15
2.2 Study of Modern Static Analysis Frameworks and AST Parsers 16
2.2.1 Tree-sitter and Concrete Syntax Trees 16
2.2.2 Code Clone Detection Paradigms 16
2.3 Limitations of Existing Systems 17
2.4 Summary of Identified Research Gaps 18

3 Requirement Analysis **19**

3.1 Introduction 19
3.2 Hardware Requirements 19
3.3 Software Requirements 20
3.4 Functional and Non-Functional Requirements 20
3.4.1 Functional Requirements 20
3.4.2 Non-Functional Requirements 22
3.5 Feasibility Study 23
3.6 Technology Stack 23

4 System Design **25**
4.1 System Architecture 25
4.2 Data Flow Diagram (DFD) 26
4.2.1 DFD Level 0 (Context Diagram) 27
4.2.2 DFD Level 1 27
4.2.3 DFD Level 2 (Inference / Debt Scoring Process Decomposition) . . 28
4.3 Activity Diagram 29

5 Implementation **31**
5.1 Introduction 31
5.2 System Algorithms 32
5.2.1 AST Parsing and Complexity Extraction 32
5.2.2 Sliding-Window Token Hashing and Clone Detection 33
5.2.3 Explainable Technical Debt Scoring 34
5.3 Implementation of the Core Analysis Engine and AST Parsing 34
5.3.1 Lexical Scanning and Line Counting 34
5.3.2 Tree-sitter Grammar Integration and AST Traversal 34
5.4 Technical Debt Scoring Engine and Explainable Risk Heuristics 35
5.5 Implementation of Exact Token Clone Detection 35
5.6 In-Memory Storage, Asynchronous Scan Queue, and Developer Dashboard 37
5.6.1 Zero-Setup In-Memory Persistence 37
5.6.2 Asynchronous Scan Queue Management 37
5.6.3 Developer-Focused React Dashboard 37

6 Testing and Results **39**
6.1 Testing Strategies and Methodologies 39
6.2 System Testing 39
6.2.1 White Box Testing 39
6.2.2 Black Box Testing 40
6.2.3 Edge Case Testing 40
6.3 Test Case Design for System Robustness 41
6.4 Runtime Performance and Throughput Benchmarking 41
6.5 Complexity and Risk Metric Evaluation 42
6.6 Duplicate Code Detection Performance 43
6.7 Comparative Analysis with Existing Tools 43
6.8 Hotspot Detection and Severity Distribution 45

7 Conclusion and Future Scope **46**
7.1 Summary 46
7.2 Proposed System Review 46
7.3 Achievement of Objectives 47

7.4 System Advantages 47
7.5 System Limitations 48
7.6 Future Direction 49

Bibliography **50**

List of Figures

4.1 System Architecture Flow of CodeInsight 25
4.2 DFD Level 0 (Context Diagram) 27
4.3 Data Flow Diagram (DFD Level 1) of the System 27
4.4 DFD Level 2 (Decomposition of Technical Debt Scoring Process) 28
4.5 Activity Diagram of the Scan and Analysis Workflow 30

6.1 Technical Debt Severity Distribution across Analyzed Test Repository . . 45

List of Tables

1.1 Project Implementation Timeline 14

3.1 CodeInsight Technology Stack 24

5.1 Technical Debt Scoring Thresholds and Weight Distributions 36

6.1 Runtime Performance and Throughput Benchmarking 42
6.2 Complexity and Nesting Depth Metrics on Test Fixture Suite 42
6.3 Duplicate Clone Detection Verification on Controlled Clones 43
6.4 Comparative Analysis with Existing Static Analysis Solutions 44

Chapter 1

Introduction

1.1 Introduction

Modern software development moves at an unprecedented pace. Development teams are constantly pushed to ship features rapidly, iterate on feedback, and release frequent updates. In this race to meet deadlines, code quality often takes a backseat to immediate functional requirements. Developers take shortcuts, write quick workarounds, copy-paste existing code snippets rather than refactoring shared utilities [7], and introduce nested conditional logic to handle new edge cases. Over time, these small compromises accumulate into what software engineering literature defines as *technical debt* [2, 6].

Just like financial debt, technical debt incurs continuous interest. As a codebase degrades, simple feature requests take significantly longer to implement, routine bug fixes introduce regressions in seemingly unrelated modules, onboarding new engineers becomes frustrating, and developer morale drops. While most engineers intuitively know which parts of their codebase are messy, intuitive impressions lack quantitative rigor. Engineering managers and tech leads cannot easily prioritize refactoring work during sprint planning without clear, objective metrics that prove where the worst risks lie.

To solve this problem, this project presents **CodeInsight**: an automated codebase health analysis and technical debt prioritization platform. CodeInsight provides an end-to-end static analysis pipeline that evaluates software repositories without requiring compilation or execution. It combines high-speed abstract syntax tree (AST) parsing, normalized token-based code clone detection, and an explainable technical debt risk scoring engine. By pairing a low-overhead analysis core with an interactive, developer-focused React dashboard, CodeInsight transforms raw static analysis telemetry into clear, prioritized refactoring targets.

1.2 Motivation

The inspiration behind CodeInsight came from observing the practical difficulties development teams experience when trying to maintain code quality. While working on

various software projects during our computer science coursework at IIIT Manipur, we noticed a recurring pattern: as repositories grow past a few thousand lines of code, certain files quietly turn into unmaintainable monsters. Functions swell to hundreds of lines with four or five levels of nested loops and conditionals, and identical logic gets copy-pasted across multiple files because nobody wants to risk touching a shared component.

When developers look for static analysis tools to address these issues, they usually encounter two extremes. On one side are lightweight linters like ESLint [11] or Clang-Format, which catch basic syntax errors and style inconsistencies but provide no high-level architectural insight into technical debt. On the other side are heavy enterprise platforms like SonarQube [10], which require complex infrastructure, dedicated database installations, substantial server RAM, and lengthy continuous integration setup before a developer can inspect a single scan.

We saw a clear opportunity to build a modern developer tool that bridges this gap: a tool that is fast, completely self-contained, runs with zero external database configuration, and provides immediate, actionable feedback on where technical debt is concentrated. Our motivation is to give individual developers and small engineering teams a transparent, understandable tool that helps them keep their codebases healthy before technical debt spirals out of control.

1.3 Problem Statement

Managing software maintainability through manual code reviews or existing static analysis platforms faces several well-documented challenges:

- **High Setup and Infrastructure Overhead:** Traditional technical debt platforms require external relational databases, container services, and administrative configuration before any analysis can run. This heavy footprint discourages adoption in small teams and local development workflows.
- **Opaque and Arbitrary Metrics:** Many existing tools present overall "maintainability ratings" (such as letter grades A through F) without clearly showing developers how those ratings were computed. When developers cannot see the mathematical formula behind a score, they distrust the recommendations.
- **Cognitive Complexity Blindness:** Simple line counts and basic style rules fail to capture the true difficulty of understanding code. A 50-line function with multiple nested loops, early returns, and branching conditions is far harder to maintain and test than a 200-line declarative configuration or lookup table.
- **Proliferation of Duplicate Code (Clones):** Copy-pasting code is one of the most common ways bugs replicate across a codebase. Standard text diffing tools

cannot detect clones when comments or whitespace vary slightly, while complex tree-to-tree comparison algorithms are too computationally expensive for large projects.

- **Lack of Actionable Remediation Triage:** Most analysis tools dump hundreds of isolated warnings onto the developer without indicating which files pose the highest risk to the project. Without clear prioritization, engineering teams suffer from alert fatigue and end up ignoring static analysis results altogether.

1.4 Objective

The primary objective of this project is to design, implement, and validate an automated, developer-focused static code analysis platform that accurately quantifies codebase health and prioritizes technical debt. The specific project objectives are:

- To build a high-performance analysis engine that extracts abstract syntax trees (AST) across multiple languages (C, C++, JavaScript, TypeScript) using Tree-sitter grammars [5].
- To implement precise code complexity metrics, including McCabe Cyclomatic Complexity ($M = 1 + D$) [1], maximum control-flow nesting depth, and function signature parameters [8].
- To design a fast, normalized sliding-window token hashing algorithm capable of detecting Type-1 exact code clones with maximal expansion while filtering out trivial boilerplate code.
- To formulate a transparent, explainable technical debt risk scoring formula ($R = 0.45C + 0.35D + 0.20S$) that breaks down risks into complexity, duplication, and structural factors.
- To implement a lightweight, zero-dependency in-memory storage layer and asynchronous scan queue that eliminates the need for external database setup.
- To create a responsive, developer-first React dashboard that visualizes technical debt distribution, highlights critical hotspots, and provides granular file-level inspections.

1.5 What is our Contribution?

In this project, I have developed CodeInsight as an integrated static analysis platform that transforms how developers inspect and prioritize code maintainability. My specific contributions include:

1. **High-Speed AST-Based Metric Extraction:** I integrated Tree-sitter parsers into a standalone analysis core to compute exact cyclomatic complexity and nesting depth directly from syntax trees, ensuring language-accurate measurement without relying on brittle regular expressions. An integrated JavaScript engine was also built as a seamless cross-platform fallback.

2. **Maximal Duplicate Clone Detection Pipeline:** I designed a token normalization and sliding-window hashing mechanism that accurately groups duplicate code fragments across files, automatically expanding matching lines to find maximal clone boundaries while suppressing false positives from repeated closing braces or short boilerplate.

3. **Explainable Technical Debt Formulation:** Rather than using a black-box rating, I established an intuitive mathematical formulation: $R = 0.45C + 0.35D + 0.20S$, where structural debt is further decomposed into physical length, nesting depth, and parameter count ($S = 0.40L + 0.35N + 0.25P$). The platform generates explicit human-readable risk factors for every file, explaining exactly why it received its score.

4. **Remediation Priority Heuristic ($P1$–$P4$):** I formulated a triage heuristic that weighs risk against estimated refactoring effort, categorizing files into actionable priority tiers so teams immediately know which files require urgent attention.

5. **High-Density Developer Dashboard:** I built an interactive React/Vite dashboard [14] inspired by tools like GitHub Insights and Linear, featuring real-time status polling, interactive distribution charts, sortable file tables, and deep modal inspectors for function diagnostics and clone clusters.

1.6 Project Timeline

The project was executed over a three-month development lifecycle from August 2026 to October 2026. The milestone breakdown across the project duration is summarized in Table 1.1.

Table 1.1: Project Implementation Timeline

Month	Activities
August 2026	Foundation and Architecture. Literature review on static code analysis, AST parsing, and technical debt models. Environment setup, C++ CLI project scaffolding with CMake, Tree-sitter integration for C/C++ and JavaScript grammars, and initial file discovery scanner.
September 2026	Core Engine Implementation. Implementing McCabe cyclomatic complexity, nesting depth calculation, and formal parameter counting. Building the sliding-window token hashing pipeline and maximal clone expansion algorithm. Developing the unified technical debt scoring formula ($R = 0.45C + 0.35D + 0.20S$) and unit testing against controlled test fixtures.
October 2026	Integration, UI, and Final Evaluation. Designing the Express REST API and background scan queue. Building the zero-setup in-memory persistence service. Developing the interactive React/Vite frontend dashboard with Recharts visualizations. Comprehensive white-box, black-box, and edge-case testing, performance benchmarking on real-world repositories, and final report documentation.

Chapter 2

Literature Survey

2.1 Existing Systems and Methodologies

Static program analysis has been a foundational pillar of software engineering for over four decades. The primary objective of static analysis is to evaluate computer software without executing the programs, enabling the identification of design flaws, security vulnerabilities, coding standard violations, and structural degradation early in the software development lifecycle. Foundational methodologies for quantifying program structure and maintainability trace back to McCabe's graph-theoretic cyclomatic complexity metric [1] and Halstead's software science formulation [9], which pioneered algorithmic analysis of source code volume and cognitive difficulty.

Early approaches to automated code evaluation relied heavily on lexical analysis and regular expressions. Tools such as Unix `grep`, `lint` for C programs introduced by Johnson in 1978 [13], and traditional style checkers scanned source code as raw streams of characters. While effective for simple pattern matching—such as identifying prohibited variable names or missing semicolons—lexical analysis suffers from severe semantic blindness. It cannot reliably distinguish between a variable declaration, a function invocation, or an occurrence within a commented block or string literal.

To overcome the limitations of lexical pattern matching, compiler frontends began providing access to compiler intermediate representations and Abstract Syntax Trees (ASTs). Modern compilers like Clang and GCC expose programmatic APIs (such as Clang LibTooling and Clang-Tidy [12]) that allow developers to write custom AST matchers and diagnostic rules. Similarly, in the interpreted and dynamic language domain, JavaScript-focused linters like JSHint and later ESLint [11] popularized modular, pluggable rule systems operating on standard abstract syntax formats like the ESTree specification.

In the commercial and enterprise domain, comprehensive software quality platforms emerged, most notably SonarQube [10] (developed by SonarSource). SonarQube aggregates multi-language analysis, test coverage tracking, rule compliance, and high-level maintainability ratings into centralized web dashboards. Similarly, tools like Code Climate and Coverity established continuous inspection pipelines inte-

grated into modern code hosting environments. These platforms shifted the industry perspective from isolated compiler warnings to holistic codebase health monitoring.

2.2 Study of Modern Static Analysis Frameworks and AST Parsers

The accuracy and responsiveness of any static analysis system depend fundamentally on how it parses and represents source code. In recent years, significant advancements in formal parsing theory and compiler construction have reshaped static analysis architectures.

2.2.1 Tree-sitter and Concrete Syntax Trees

A major breakthrough in practical code analysis was the development of **Tree-sitter** [5], an incremental parsing library originally designed by Max Brunsfeld for text editors and developer tools. Unlike traditional compiler generators (such as Yacc, Bison, or ANTLR) that require complete, syntactically flawless source code and often produce high memory overhead, Tree-sitter is built on generalized LR (GLR) parsing algorithms. It generates a concrete syntax tree (CST) that preserves every token in the original file, including whitespace, comments, and punctuation, while being exceptionally resilient to syntax errors.

Tree-sitter provides unified C-language bindings and pre-compiled grammars for dozens of programming languages. Because it builds the parse tree in a single fast pass directly into memory, it can parse thousands of lines of code in a few milliseconds. In the context of multi-language static analysis, Tree-sitter allows an analysis engine to navigate source structures across C, C++, JavaScript, TypeScript, and Python through a consistent, high-speed procedural API without needing to install individual language compiler toolchains.

2.2.2 Code Clone Detection Paradigms

Code duplication—frequently referred to in software engineering research as *code clones*—is recognized as one of the most detrimental contributors to technical debt [2, 6], famously cataloged by Fowler [7] as a primary code smell requiring disciplined refactoring. Roy and Cordy [4] have established a widely accepted taxonomy that categorizes code clones into four primary classes:

1. **Type-1 (Exact Clones):** Identical source code fragments except for minor variations in formatting, indentation, and comments.

2. **Type-2 (Parameterized / Renamed Clones):** Syntactically identical fragments where identifiers, variable names, type names, and literal values have been altered.

3. **Type-3 (Near-Miss Clones):** Copied fragments with added, deleted, or reordered statements.

4. **Type-4 (Semantic Clones):** Computationally and semantically equivalent code blocks that perform the identical functional logic but utilize entirely different syntactic constructs and algorithms.

Extensive empirical evaluations, such as the landmark tool comparison study conducted by Bellon et al. [3], demonstrated that while AST-based and graph-based clone detectors (such as CloneDR) can identify complex Type-3 and Type-4 clones, they suffer from high algorithmic complexity ($O(N^2)$ or worse) and significant memory footprints. In contrast, token-based and line-based sliding-window approaches (such as CCFinder by Kamiya et al.) achieve remarkable computational efficiency, making them capable of scanning large-scale production codebases within seconds while maintaining near-perfect precision for Type-1 and normalized clones.

2.3 Limitations of Existing Systems

Despite the widespread availability of static analysis tools, development teams frequently struggle to integrate them effectively into their daily engineering routines. Our review of existing systems identified several critical limitations:

- **Heavy Infrastructure and Database Dependency:** Enterprise platforms such as SonarQube [10] require dedicated relational databases (e.g., PostgreSQL, Oracle), containerization runtimes, and substantial background server processes. For individual developers, open-source maintainers, or small development teams seeking quick feedback on a local project, setting up and maintaining this infrastructure represents an unnecessary barrier to entry.

- **Black-Box Scoring and Arbitrary Metrics:** Many commercial tools assign arbitrary quality grades (such as “Maintainability Rating A” or “Technical Debt: 3 days”) based on proprietary internal formulas. When developers cannot trace how a specific number or letter grade was calculated, they view the tool with skepticism and struggle to justify refactoring tasks to engineering leadership.

- **Information Overload and Alert Fatigue:** Traditional linters report thousands of individual stylistic infractions—ranging from trailing whitespace to arbitrary naming conventions—without weighting them by their structural risk.

Developers become overwhelmed by noisy warning lists and frequently end up disabling the tool altogether.

- **Language Fragmentation:** Developers working in multi-language or full-stack environments (for example, combining C++ native modules with JavaScript/TypeScript services) are typically forced to configure multiple disjoint static analysis tools, each with its own output format, configuration file, and reporting structure.

2.4 Summary of Identified Research Gaps

Our comprehensive literature survey reveals a distinct research and tooling gap between lightweight, shallow linters and heavyweight enterprise platforms. Specifically:

1. There is a strong need for an **independent, zero-setup static analysis engine** that operates with self-contained in-memory persistence, allowing instantaneous repository scans without external database dependencies.

2. There is a demand for a **transparent and mathematically explainable technical debt model** that integrates cyclomatic complexity, code duplication, and structural nesting into a clear, weighted formulation ($R = 0.45C + 0.35D + 0.20S$) adhering to modern clean code craftsmanship principles [8] accompanied by explicit human-readable risk factors.

3. There is a lack of developer tools that combine **fast AST-level function diagnostics** with **clustered clone visualization** within a high-density, interactive user interface designed specifically for software engineers.

By directly addressing these shortcomings, the CodeInsight platform was designed and built to provide a fast, accessible, and mathematically rigorous solution for codebase health monitoring.

Chapter 3

Requirement Analysis

3.1 Introduction

The requirement analysis phase defines the functional capabilities, operational constraints, and engineering specifications necessary to build the CodeInsight platform. The goal of CodeInsight is to provide software engineers with an immediate, accurate, and completely self-contained static analysis solution.

Unlike conventional enterprise code quality tools that impose heavy setup burdens, CodeInsight was conceived from the ground up to operate with zero external database dependencies, providing instantaneous codebase health evaluations through a clean single-server or development setup. Documenting both functional and non-functional requirements guarantees that the implemented platform is computationally efficient, deterministic, user-friendly, and maintainable.

3.2 Hardware Requirements

One of the primary design goals of CodeInsight is to ensure accessibility on standard developer workstations, avoiding the high memory and CPU overhead associated with enterprise analysis suites. The hardware requirements for running and developing CodeInsight are outlined below:

- **Central Processing Unit (CPU):** Dual-core or quad-core modern processor (Intel Core i3/i5/i7, AMD Ryzen 3/5/7, or Apple Silicon). The multi-threaded C++ analysis engine and asynchronous Node.js event loop leverage multi-core architectures during large-scale directory scanning and token hashing.
- **System Memory (RAM):** A minimum of **4 GB** of RAM is required, with **8 GB** recommended for analyzing repositories containing over 100,000 lines of code. Because CodeInsight utilizes an efficient in-memory data store rather than running a heavy Java Virtual Machine or relational database server, its memory footprint remains well below 150 MB during normal operation.
- **Storage Space:** Approximately **250 MB** of available disk storage for the codebase, compiled native binaries, and node dependencies. No dedicated database

disk storage allocation is needed.

- **Display Resolution:** A minimum screen resolution of 1280×720 pixels, with 1920×1080 (Full HD) recommended to comfortably view the high-density metrics tables, distribution charts, and side-by-side code clone inspections.
- **Network:** Internet connection for installing dependencies and development tools.

3.3 Software Requirements

The software environment for CodeInsight was selected to ensure maximum cross-platform portability across Windows, Linux, and macOS environments:

- **Operating System:** Windows 10/11 (64-bit), Ubuntu Linux 20.04+ (or equivalent POSIX Linux distribution), or macOS 12+.
- **Runtime Environment: Node.js** version 18.0 or higher (actively tested and validated on Node.js v24.x LTS) with the **npm** package manager.
- **C++ Compilation Toolchain (for Engine Development):** A modern C++17 compatible compiler (GCC 9+, Clang 10+, or MSVC 2019+) along with **CMake 3.15+** for building the standalone native analyzer executable.
- **Tree-sitter Parsing Grammars:** Tree-sitter C runtime library [5] along with pre-compiled grammars for C, C++, and JavaScript.
- **Web Browser:** Any modern evergreen web browser supporting ECMAScript 2020 and CSS Grid (Google Chrome 90+, Mozilla Firefox 88+, Microsoft Edge 90+, or Apple Safari 14+).
- **Development and Build Tools:** Visual Studio Code (or equivalent IDE) and Vite 5 [14] for fast frontend hot module replacement and optimized production asset bundling.

3.4 Functional and Non-Functional Requirements

3.4.1 Functional Requirements

The functional requirements describe the specific operations and capabilities that the system must provide:

1. **Repository Selection:** The system shall allow users to specify or select a local source code repository for analysis. It shall validate the repository path and identify supported source files before beginning the scanning process.

2. **Source Code Scanning:** The system shall recursively scan the selected repository and identify supported source code files. It shall collect the information required for subsequent analysis while excluding irrelevant directories and unsupported files according to the configured rules.

3. **Cyclomatic Complexity Analysis:** The system shall analyze supported source files to identify functions, decision points, and nesting depth. It shall calculate cyclomatic complexity and convert the extracted metrics into a normalized complexity score between 0 and 100.

4. **Code Duplication Detection:** The system shall identify repeated code fragments using the implemented duplication detection technique. It shall calculate the percentage of duplicated source lines and convert the result into a normalized duplication score.

5. **Structural Analysis:** The system shall calculate structural metrics, including lines of code and unusually long lines. These measurements shall be used to determine a normalized structural score between 0 and 100.

6. **Technical Debt Risk Calculation:** The system shall combine the normalized complexity, duplication, and structural scores using configurable weights. The active weights shall be normalized to account for analysis dimensions that are not implemented in the current version. The overall risk score shall remain within the range of 0 to 100.

7. **Risk Classification:** The system shall classify analyzed files into four risk categories based on their overall risk scores:

   - **Low:** 0–24
   - **Moderate:** 25–49
   - **High:** 50–74
   - **Critical:** 75–100

   These categories represent heuristic estimates of potential maintenance concerns rather than the probability of software defects.

8. **Results Visualization:** The system shall display analysis results through a graphical user interface. The dashboard shall present repository-level summaries, risk distributions, file-level scores, and the relative contribution of individual metrics.

9. **File-Level Inspection:** The system shall allow users to inspect individual files and view their complexity, duplication, structural scores, overall risk score, and relevant metric-based explanations.

10. **Filtering and Sorting:** The system shall provide options to filter and sort analyzed files according to risk category, risk score, file path, or other available metrics.

11. **Report Export:** The system should allow users to export analysis results in a structured format, such as JSON or CSV, so that results can be stored and reviewed outside the application.

12. **Error Handling:** The system shall handle invalid repository paths, unreadable files, unsupported source formats, parsing failures, and analyzer execution errors without terminating the entire application unnecessarily.

3.4.2 Non-Functional Requirements

The non-functional requirements govern the quality attributes, operational constraints, and architecture of the platform:

1. **Performance:** The system should analyze repositories within a reasonable time based on repository size, file count, and source code complexity. File processing and metric calculation should be implemented efficiently to avoid unnecessary computational overhead.

2. **Accuracy and Consistency:** The system shall apply consistent metric definitions and scoring rules across analyzed files. The same source code and configuration should produce reproducible results, subject to the deterministic behavior of the implemented analysis tools.

3. **Usability:** The user interface shall provide clear navigation, readable tables, meaningful visualizations, and understandable explanations of risk scores. Users should be able to identify high-risk files without requiring extensive knowledge of the internal scoring implementation.

4. **Modularity:** The system shall separate the frontend, backend, code analysis engine, and risk-scoring logic into distinct components. This modular structure shall support testing, maintenance, and future enhancements.

5. **Maintainability:** The implementation shall use readable code, consistent naming conventions, reusable modules, and appropriate error handling. Scoring thresholds and weights should be configurable where practical.

6. **Reliability:** The system shall handle common input and processing errors gracefully. A failure involving an individual source file should not unnecessarily prevent the analysis of other readable files.

7. **Security:** The system shall validate repository paths, restrict file access to the intended analysis scope, and handle external process execution safely. It shall avoid unsafe shell-command construction and protect against path traversal and unintended file modification.

8. **Scalability:** The system architecture should permit future improvements such as incremental analysis, background scanning, persistent scan history, dependency analysis, and Git-based metrics.

3.5 Feasibility Study

A thorough feasibility evaluation was conducted before and during development:

- **Technical Feasibility:** Utilizing Tree-sitter provides production-grade concrete syntax trees with exceptional fault tolerance, allowing incomplete or non-compiling source files to be parsed cleanly. Pairing a high-speed C++ engine with an Express API and a React frontend combines low-level algorithmic efficiency with rapid web development. The integration proved robust and feasible across platforms.

- **Operational Feasibility:** Traditional static analysis tools often fail in practice because developers find database configuration and maintenance too cumbersome. By decoupling CodeInsight from external databases like PostgreSQL and utilizing a fast in-memory storage layer, operational friction was completely eliminated. Any developer can clone the repository, run `npm start`, and begin analyzing repositories immediately.

- **Economic Feasibility:** The entire technology stack consists of open-source, permissively licensed libraries (MIT, Apache 2.0). CodeInsight incurs zero commercial licensing fees, requires no paid third-party cloud APIs, and can be hosted on minimal compute infrastructure, making it highly cost-effective for academic and commercial use.

3.6 Technology Stack

Table 3.1 summarizes the technologies and frameworks selected to build the CodeInsight static analysis platform.

Table 3.1: CodeInsight Technology Stack

Category	Technology / Library	Purpose in System
Analysis Engine	C++17, GCC / Clang / MSVC	Native high-speed scanner and AST parsing executable
Parsing Library	Tree-sitter	Concrete syntax tree generation for C, C++, and JavaScript
Build System	CMake 3.15+	Cross-platform native binary compilation
Backend Runtime	Node.js (v18+ / v24 LTS)	Orchestration server, asynchronous scan worker, REST API
Backend Framework	Express.js 4.x	REST routing, request handling, and static asset serving
Persistence Layer	In-Memory Store (`StorageService`)	Zero-dependency in-memory store for repos, scans, and metrics
Frontend Framework	React 18	Declarative component-based user interface
Frontend Tooling	Vite 5	Rapid development server (HMR) and optimized Rollup bundling
Data Visualization	Recharts	Interactive technical debt distribution bar charts
Styling	Custom CSS Tokens	Dark theme palette, monospaced typography, severity badges
Testing Suites	Custom Node.js Test Harness	Automated unit and integration testing across all phases

Chapter 4

System Design

4.1 System Architecture

The system architecture of CodeInsight is designed as an end-to-end, sequential processing pipeline that transforms raw source code files into structured, actionable technical debt telemetry. Rather than relying on a heavy monolithic architecture or external database servers, the platform separates user interaction, API orchestration, AST parsing, and risk evaluation into distinct, well-coordinated stages. Figure 4.1 illustrates the sequential architecture flow of the platform from initial repository selection to final visualization.

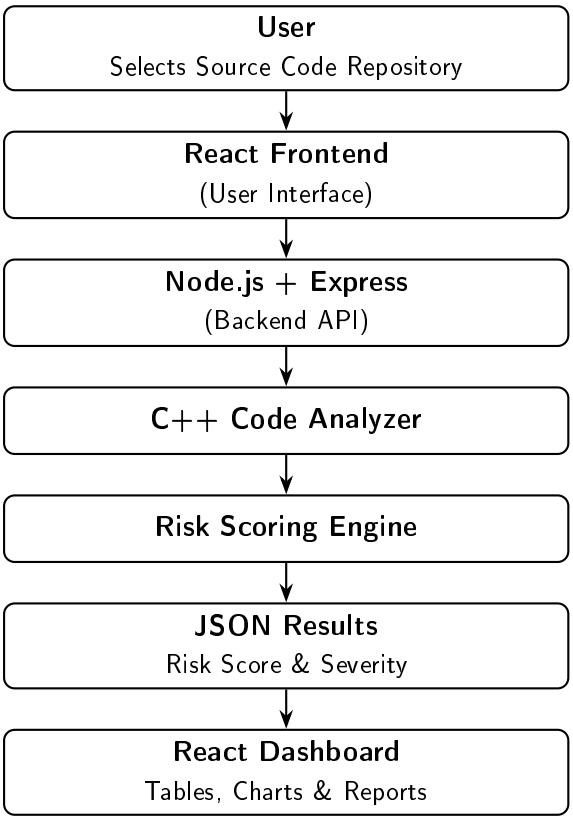


Figure 4.1: System Architecture Flow of CodeInsight

The pipeline executes through seven well-defined stages, ensuring deterministic data transformation at each phase:

1. **User (Selects Source Code Repository):** The workflow begins with the user selecting or specifying the target source code repository on the local filesystem. The user provides the directory path to be evaluated through the web interface.

2. **React Frontend (User Interface):** The React user interface captures the user's input, validates the path format on the client side, and initiates an asynchronous scan request by dispatching a REST call to the backend server.

3. **Node.js + Express (Backend API):** The Node.js Express server acts as the API gateway and orchestration layer. It manages the asynchronous scan queue, validates directory existence and permissions on the host system, and spawns the native analysis engine process while tracking scan execution state.

4. **C++ Code Analyzer:** The high-performance C++ engine performs recursive directory traversal with intelligent file filtering. It parses supported source files into concrete syntax trees using Tree-sitter grammars [5], computes exact McCabe cyclomatic complexity ($M = 1 + D$) [1] and maximum control-flow nesting depth, and detects Type-1 code clones through a normalized sliding-window token hashing pipeline.

5. **Risk Scoring Engine:** The risk scoring module processes the extracted metrics (complexity, duplication, and structural dimensions). It computes a normalized composite technical debt risk score ($R = 0.45C + 0.35D + 0.20S$) and assigns clear remediation priority tiers ($P1$–$P4$) along with human-readable risk factors.

6. **JSON Results (Risk Score & Severity):** The aggregated metrics, file-level diagnostics, clone clusters, and severity rankings are packaged into structured JSON payloads. This intermediate representation enables zero-dependency in-memory persistence and standardized export serialization.

7. **React Dashboard (Tables, Charts & Reports):** The frontend ingests the structured JSON telemetry, rendering an interactive developer dashboard complete with repository-level summary cards, debt distribution charts, searchable and sortable file tables, detailed function inspection modals, and export capabilities (JSON, CSV, Markdown).

4.2 Data Flow Diagram (DFD)

Data Flow Diagrams provide a visual mapping of how information enters, flows through, and exits the CodeInsight platform. They model the functional transformations applied to source files from raw repository input to completed health reports.

4.2.1 DFD Level 0 (Context Diagram)

The Context Diagram (Level 0) represents the highest-level operational boundary of CodeInsight, viewing the entire platform as a single centralized system interacting with the software engineer and the target file system.

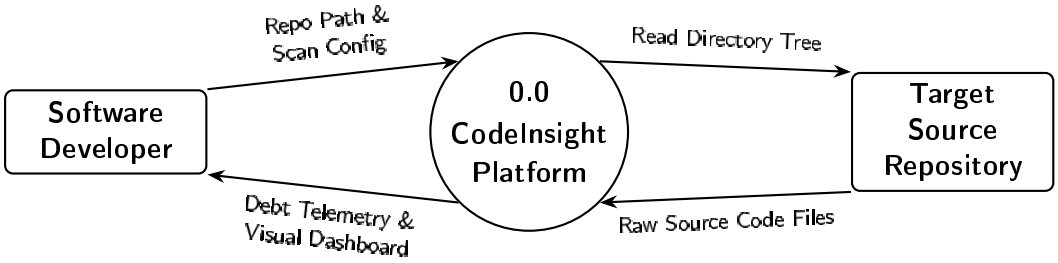


Figure 4.2: DFD Level 0 (Context Diagram)

As depicted in Figure 4.2, the developer supplies a local file system directory path and optional risk formula weights. The platform accesses the target codebase, reads raw source code files, executes parsing and scoring, and returns comprehensive debt telemetry, clone clusters, and visual summaries to the user interface.

4.2.2 DFD Level 1

The Level 1 Data Flow Diagram decomposes the system into its primary operational sub-processes and tracks data movement into the in-memory data store, as shown in Figure 4.3.

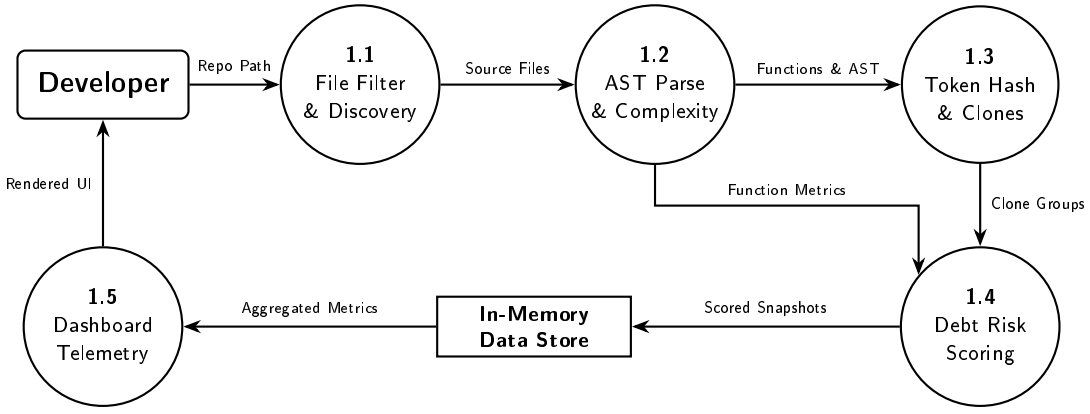


Figure 4.3: Data Flow Diagram (DFD Level 1) of the System

In Figure 4.3, the repository path passes to Process 1.1 (File Filtering & Discovery), which ignores minified files, vendor dependencies, and binary assets. Valid source files flow into Process 1.2 (AST Parsing & Complexity Extraction) to determine function-level decision points and nesting depths. Concurrently, Process 1.3 (Token Hashing & Clones) executes sliding-window matching to identify duplicate blocks. Both metric

streams feed into Process 1.4 (Debt Risk Scoring), which computes the composite risk formula and populates the In-Memory Data Store. Finally, Process 1.5 (Dashboard Telemetry) pulls the stored snapshots and renders real-time visual reports.

4.2.3 DFD Level 2 (Inference / Debt Scoring Process Decomposition)

Figure 4.4 provides a granular decomposition of Process 1.4 (Debt Risk Scoring), illustrating the mathematical steps and data transformations required to evaluate technical debt and assign severity tiers.

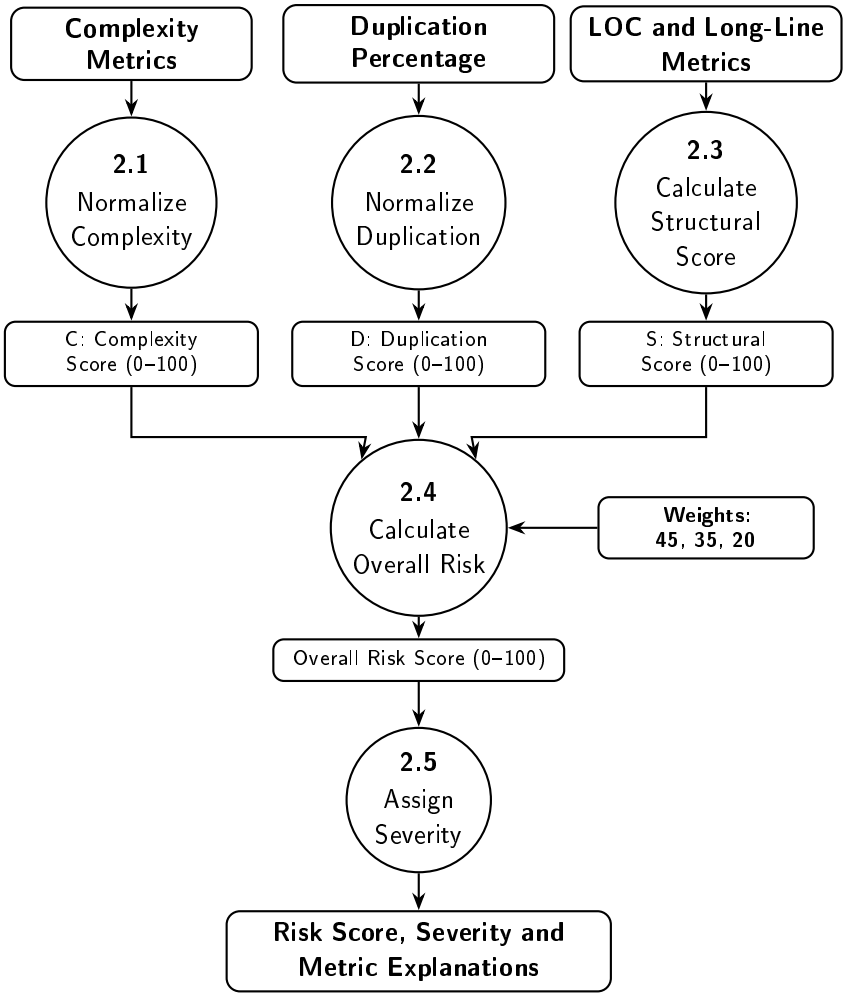


Figure 4.4: DFD Level 2 (Decomposition of Technical Debt Scoring Process)

As illustrated in Figure 4.4, the technical debt evaluation pipeline decomposes into five distinct sub-processes:

- **Sub-process 2.1 (Normalize Complexity):** Ingests raw function complexity metrics (McCabe cyclomatic complexity $M = 1 + D$) and scales them into a normalized complexity score $C \in [0, 100]$.

- **Sub-process 2.2 (Normalize Duplication):** Evaluates multi-file code clone percentages identified by the token hashing pipeline, producing a normalized duplication score $D \in [0, 100]$.
- **Sub-process 2.3 (Calculate Structural Score):** Aggregates physical line metrics, unusually long lines, and control-flow nesting depth to derive a composite structural health score $S \in [0, 100]$.
- **Sub-process 2.4 (Calculate Overall Risk):** Synthesizes the normalized scores (C, D, S) with calibrated weights (45%, 35%, and 20% respectively) according to the formulation $R = 0.45C + 0.35D + 0.20S$, producing an overall risk score bound within $[0, 100]$.
- **Sub-process 2.5 (Assign Severity):** Evaluates the overall risk score against discrete heuristic thresholds, categorizing files into Low, Moderate, High, or Critical severity tiers and generating actionable, metric-based remediation explanations.

4.3 Activity Diagram

The Activity Diagram in Figure 4.5 illustrates the operational workflow executed when a developer triggers a repository scan through the CodeInsight interface.

As shown in Figure 4.5, execution begins when a developer specifies a repository path in the UI. If the path does not exist on disk, an immediate notification is displayed and input focus returns to the prompt. Once verified, the worker thread filters irrelevant assets, computes AST-based cyclomatic complexity and nesting levels, hashes normalized token windows to uncover maximal clone blocks, calculates the composite technical debt risk formula, saves the snapshot to in-memory storage, and renders the updated dashboard views.

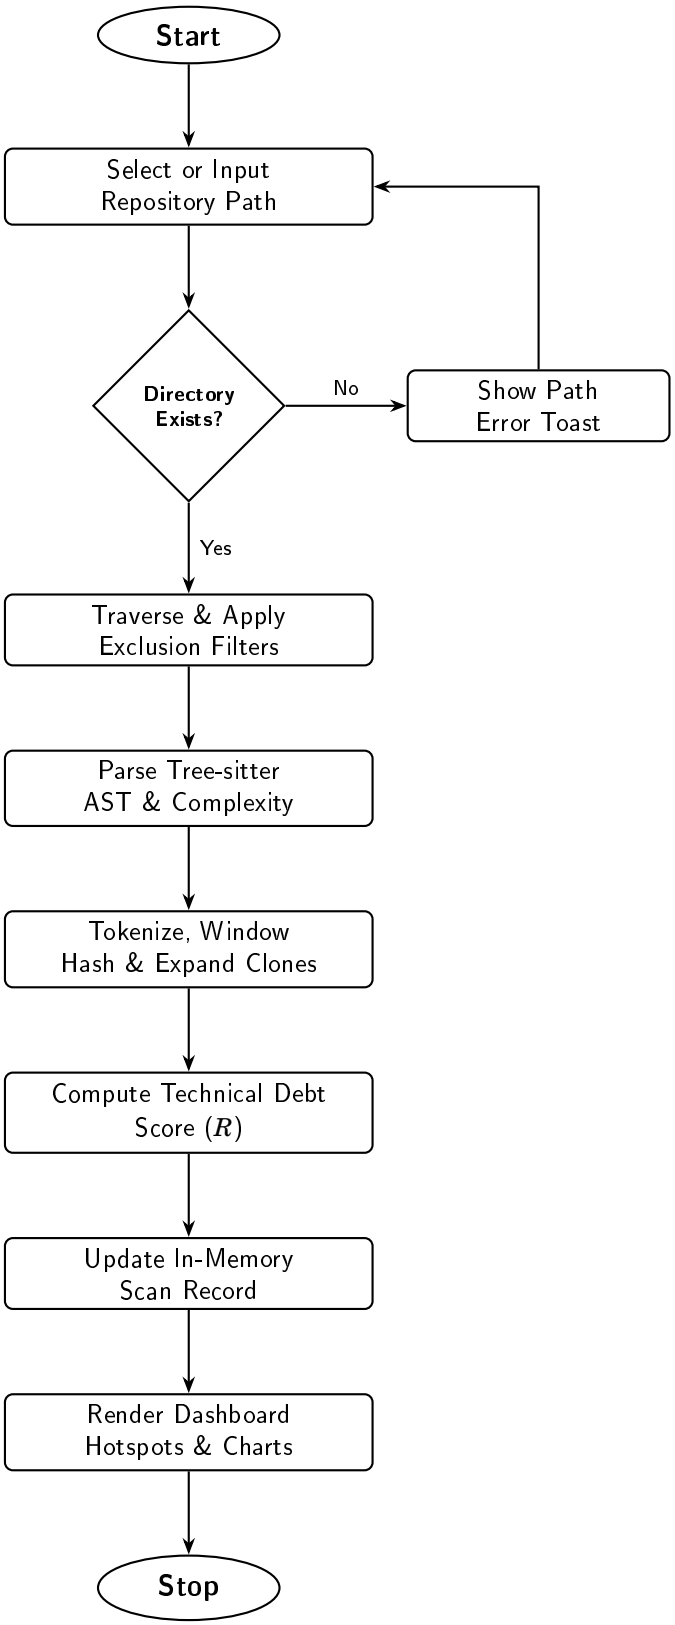


Figure 4.5: Activity Diagram of the Scan and Analysis Workflow

Chapter 5

Implementation

5.1 Introduction

Implementation is the phase in which the proposed system architecture and design are converted into a working software application. This chapter describes the implementation of CodeInsight, an automated codebase health and technical debt prioritization platform.

The system is developed using a modular architecture consisting of a React-based frontend, a Node.js and Express backend, and a C++ analysis engine. The frontend provides the user interface for repository selection and result visualization, while the backend coordinates repository scanning and communication with the analysis engine. The C++ engine performs source code analysis and extracts relevant metrics.

The extracted metrics are processed by the risk-scoring module to calculate complexity, duplication, and structural scores. These scores are combined using a weighted scoring model to estimate the relative maintenance risk of individual source files. The resulting information is returned to the frontend in JSON format and presented through tables, charts, and file-level details.

The current implementation focuses on repository scanning, static code analysis, duplication detection, structural analysis, risk calculation, and result visualization. Database integration, dependency coupling analysis, and Git modification history analysis are considered future enhancements.

5.2 System Algorithms

The core operational logic of CodeInsight is driven by three primary algorithms governing source code parsing, duplicate clone identification, and technical debt risk calculation.

5.2.1 AST Parsing and Complexity Extraction

Algorithm 1 outlines how source code files are parsed into abstract syntax trees to extract cyclomatic complexity [1] and nesting depths.

Algorithm 1 AST Parsing and Cyclomatic Complexity Extraction

Require: Source file path F, Language Grammar G_L
Ensure: Function metric list $\mathcal{M} = \{(f_i, M_i, N_i, P_i, L_i)\}$

Read raw file content $C \leftarrow$ READFILE(F)
Initialize Tree-sitter parser P with grammar G_L
Concrete Syntax Tree $T \leftarrow$ PARSE(P, C)
$\mathcal{M} \leftarrow \emptyset$
for each node $V \in$ TRAVERSENODES(T) **do**
    if V.type $\in$ {function_definition, method_definition, arrow_function} **then**
        $M \leftarrow 1$ ▷ Baseline sequential control flow
        $N_{\max} \leftarrow 0$, $N_{\text{curr}} \leftarrow 0$
        $P \leftarrow$ COUNTFORMALPARAMETERS(V)
        $L \leftarrow V$.endLine $- V$.startLine $+ 1$
        for each child node u in subtree rooted at V **do**
            if u.type $\in$ {if_statement, for_statement, while_statement, case_statement, ...} **then**
                $M \leftarrow M + 1$ ▷ Increment decision predicate count
            else if u.type $\in$ {&&, ||, ??} **then**
                $M \leftarrow M + 1$ ▷ Short-circuit logical operator branch
            end if
            if u enters new nested block **then**
                $N_{\text{curr}} \leftarrow N_{\text{curr}} + 1$
                $N_{\max} \leftarrow \max(N_{\max}, N_{\text{curr}})$
            else if u exits nested block **then**
                $N_{\text{curr}} \leftarrow N_{\text{curr}} - 1$
            end if
        end for
        $\mathcal{M} \leftarrow \mathcal{M} \cup \{(V$.name$, M, N_{\max}, P, L)\}$
    end if
end for
return $\mathcal{M}$

5.2.2 Sliding-Window Token Hashing and Clone Detection

Algorithm 2 details the sliding-window token hashing and maximal forward expansion technique utilized to identify Type-1 duplicate code fragments across files.

Algorithm 2 Sliding-Window Token Hashing and Maximal Clone Expansion

Require: Analyzed file set $F = \{F_1, \ldots, F_m\}$, Window Size $K = 5$, Min Token Length $\tau = 30$
Ensure: Clustered Duplicate Groups $G = \{G_1, G_2, \ldots\}$

Initialize hash table H : HashValue $\rightarrow$ List of (FileID, LineStart)
for each file $F \in F$ **do**
    Strip comments and empty lines from F, recording original line numbers L
    Normalize whitespace: collapse consecutive spaces and tabs to a single space
    for index $i = 1$ to $|F.\text{lines}| - K + 1$ **do**
        Block $B \leftarrow$ JOINLINES(F.lines$[i \ldots i + K - 1]$)
        if LENGTHWITHOUTWHITESPACE(B) $\geq \tau$ **then**
            $h \leftarrow$ HASH(B) ▷ Compute 64-bit rolling hash
            Append (F.id, i) to $H[h]$
        end if
    end for
end for
Candidate pairs $P \leftarrow \emptyset$
for each hash entry $h \in H$ where $|H[h]| \geq 2$ **do**
    Expand matching lines forward across paired fragments until next mismatch
    Merge overlapping blocks and filter out sub-slices covered by larger blocks
    Add maximal non-subsumed clone fragments to P
end for
Cluster connected pairs in P into disjoint duplicate groups G
return G

5.2.3 Explainable Technical Debt Scoring

Algorithm 3 demonstrates the calculation of file-level composite technical debt risk scores and priority tiers.

Algorithm 3 Explainable Technical Debt Scoring and Priority Triage

Require: File metrics $(M_{\max}, P_{\text{dup}}, \text{LOC}, N_{\max}, P_{\max})$, Weight vectors $(w_c, w_d, w_s) = (0.45, 0.35, 0.20)$

Ensure: Risk Score $R \in [0, 100]$, Severity Tier, Priority Rating ($P1$–$P4$)

$C \leftarrow$ MAPCOMPLEXITYSCORE($M_{\max}$) ▷ Table 5.1

$D \leftarrow$ MAPDUPLICATIONSCORE(P_{dup}) ▷ Table 5.1

$L \leftarrow$ MAPLENGTHSCORE(LOC), $N \leftarrow$ MAPNESTINGSCORE($N_{\max}$), $P \leftarrow$ MAPPARAMSCORE($P_{\max}$)

$S \leftarrow 0.40L + 0.35N + 0.25P$ ▷ Compute structural risk score

$R \leftarrow \min(100, \max(0, 0.45C + 0.35D + 0.20S))$

Severity $\leftarrow$ CATEGORIZESEVERITY(R) ▷ Low, Moderate, High, Critical

ImpactMultiplier $\leftarrow (M_{\max} > 15 \text{ or } C_{\text{high}} > 0)\ ?\ 1.3 : 1.0$

EffortFactor $\leftarrow \text{LOC} > 300\ ?\ 1.5 : 1.0$

PriorityScore $\leftarrow \min(100, (R \times \text{ImpactMultiplier})/\text{EffortFactor})$

Assign priority tier $P1, P2, P3,$ or $P4$ based on R and PriorityScore

return $(R, \text{Severity}, \text{PriorityScore})$

5.3 Implementation of the Core Analysis Engine and AST Parsing

The core analysis engine was implemented as a high-performance C++ executable (`codeinsight-analyzer`) leveraging Tree-sitter C runtime libraries [5] and official grammars for C, C++, and JavaScript.

5.3.1 Lexical Scanning and Line Counting

Before parsing ASTs, the engine performs lexical state-machine scanning to compute physical line metrics: total lines of code (LOC), source lines of code (SLOC), comment lines, and blank lines. A lexical state machine tracks multi-line comment delimiters (`/* ... */`) and single-line comments (`//`), ensuring that comment lines within string literals or macros are accurately classified.

5.3.2 Tree-sitter Grammar Integration and AST Traversal

The Tree-sitter parser compiles grammar definitions into fast state tables. For C and C++, the parser recognizes function declarations, constructor definitions, and member

methods. For JavaScript and TypeScript, the engine traverses function declarations, function expressions, object method definitions, and ES6 arrow functions.

During recursive subtree traversal, cyclomatic complexity increments by 1 whenever an AST node matches a branching decision point. In addition to standard control-flow keywords (`if`, `for`, `while`, `do`, `case`), the engine explicitly counts short-circuit boolean operators (`&&`, `||`, and nullish coalescing `??`) because each short-circuit creates a distinct branch in the program's control-flow graph. Maximum nesting depth is tracked by maintaining a stack of enclosing compound statement blocks.

5.4 Technical Debt Scoring Engine and Explainable Risk Heuristics

To eliminate the black-box nature of commercial code rating tools, CodeInsight computes technical debt using a transparent, explainable mathematical formulation:

$$\boxed{R = 0.45C + 0.35D + 0.20S} \tag{5.1}$$

where the structural risk score S is computed from physical length (L), maximum nesting depth (N), and maximum parameter count (P):

$$\boxed{S = 0.40L + 0.35N + 0.25P} \tag{5.2}$$

Table 5.1 summarizes the empirical threshold mappings used to convert raw extracted metrics into normalized sub-scores (0 to 100).

Based on the calculated risk score R, each file is classified into a clear severity tier:

- **Low Risk** ($0 \le R < 25$)**:** Clean, well-structured, maintainable code.
- **Moderate Risk** ($25 \le R < 50$)**:** Moderate complexity; monitor during regular development sprints.
- **High Risk** ($50 \le R < 75$)**:** Significant technical debt; target for refactoring.
- **Critical Risk** ($75 \le R \le 100$)**:** Severe technical debt hotspot; top priority candidate for decomposition.

5.5 Implementation of Exact Token Clone Detection

Code duplication detection is implemented using normalized sliding-window token hashing. The detector processes source files in three phases:

Table 5.1: Technical Debt Scoring Thresholds and Weight Distributions

Metric Category	Raw Input Value	Mapped Score	Formula Weight
Complexity (C)	0 (No functions / Config file)	0	45%
	$1-5$ decision points	10	
	$6-10$ decision points	35	
	$11-15$ decision points	65	
	>15 decision points	90	
Duplication (D)	0% duplicated lines	0	35%
	$1\%-5\%$ duplicated lines	20	
	$6\%-10\%$ duplicated lines	40	
	$11\%-20\%$ duplicated lines	70	
	$>20\%$ duplicated lines	90	
Structural (S)	$S = 0.40L + 0.35N + 0.25P$		20%
Length (L)	≤ 100 LOC and func lines ≤ 30	10	(40% of S)
	$101-300$ LOC or func $31-60$	35	
	$301-700$ LOC or func $61-120$	65	
	>700 LOC or func >120	90	
Nesting (N)	Max nesting depth ≤ 1	10 (0 if no func)	(35% of S)
	Max nesting depth $= 2$	35	
	Max nesting depth $= 3$	65	
	Max nesting depth ≥ 4	90	
Parameters (P)	Max parameters ≤ 2	10 (0 if no func)	(25% of S)
	Max parameters $3-4$	35	
	Max parameters $5-6$	65	
	Max parameters >6	90	

1. **Token Normalization and Source Mapping:** Comments and empty lines are removed, while maintaining a bidirectional line mapping to the original source file. Internal whitespace sequences are collapsed to single spaces.

2. **Sliding Window Rolling Hashing:** A sliding window of size $K = 5$ lines moves across normalized lines. To eliminate false-positive duplicate clusters caused by trivial syntax (such as repeated closing braces or boilerplate declarations), blocks with fewer than $\tau = 30$ non-whitespace characters are discarded. Valid blocks are hashed using a 64-bit rolling hash.

3. **Maximal Clone Expansion and Subsumption Filtering:** When identical hash keys occur across multiple file locations, the algorithm expands matching lines forward line-by-line until a mismatch occurs. A subsumption filter removes shorter fragments that are completely contained within longer clone blocks, clustering the resulting fragments into unique duplicate groups.

5.6 In-Memory Storage, Asynchronous Scan Queue, and Developer Dashboard

5.6.1 Zero-Setup In-Memory Persistence

To eliminate the installation, configuration, and maintenance overhead of external database servers such as PostgreSQL, CodeInsight uses a high-performance in-memory persistence layer implemented in `StorageService`. Repositories, scans, file snapshots, and duplicate groups are stored in optimized JavaScript `Map` structures indexed by unique identifiers. The service provides instantaneous CRUD operations ($O(1)$ lookup time) and automatically registers default test fixtures on startup.

5.6.2 Asynchronous Scan Queue Management

Repository analysis is executed asynchronously through `ScanQueueService`. When a scan request is submitted, a scan record is created with status `queued`. The queue worker transitions the job to `scanning`, invokes the analysis engine, aggregates metrics, evaluates risk scores, and marks the job as `completed`. The frontend queries the job status every 2 seconds via a lightweight status endpoint, ensuring a smooth, non-blocking user experience even on larger codebases.

5.6.3 Developer-Focused React Dashboard

The frontend interface is engineered specifically for software developers using React and Vite [14]:

- **Overview Dashboard:** Displays aggregated repository statistics (total files, lines of code, average complexity, duplicate percentage, overall risk score), an interactive Recharts risk distribution bar chart, a live formula card explaining $R = 0.45C + 0.35D + 0.20S$, and the top technical debt hotspots requiring remediation.
- **Searchable File Explorer:** Features real-time substring search, multi-select filters for language and severity, multi-column sorting (by risk score, lines, or complexity), pagination, and quick-launch modal triggers.
- **Granular File Inspector Modal (`FileDetail`):** Provides a comprehensive breakdown of an individual file's score, displaying exact normalized sub-points (C, D, S, L, N, P), human-readable risk factors, and a complete table of all functions with their cyclomatic complexity, nesting depth, and parameter counts.
- **Clustered Duplicate Clone Viewer (`DuplicateViewer`):** Allows developers to inspect duplicate code clusters, filter by minimum line count, view participating file fragments, and inspect syntax-highlighted code blocks.

Chapter 6

Testing and Results

6.1 Testing Strategies and Methodologies

Rigorous software verification is essential to ensure that static analysis tools produce accurate, deterministic, and statistically reliable metrics. If an automated code health platform reports inaccurate complexity scores or misses substantial duplicate code blocks, developers will quickly lose trust in its recommendations.

In this project, testing was conducted systematically across multiple tiers of the system architecture. Unit testing was employed during the early implementation phases to verify individual functions in isolation—such as the lexical comment stripping state machine, line normalization algorithms, and token rolling hash routines. Deterministic synthetic test files were fed into these functions, asserting that extracted token counts and hashes matched mathematically calculated expected values. Integration tests were subsequently executed to confirm that discrete modules—including the directory scanner passing file paths to the Tree-sitter parser, and the AST engine feeding function lists to the risk scorer—communicated smoothly without type errors or data loss.

6.2 System Testing

To evaluate the overall reliability and robustness of CodeInsight, comprehensive system testing was executed across three core testing paradigms: white-box testing, black-box testing, and edge-case testing.

6.2.1 White Box Testing

White-box testing examined the internal logic and control flow of the analysis engine and scoring services:

- **AST Decision Point Accounting:** Verified that every branching node type in C++ and JavaScript (such as `if`, `else if`, `for`, `while`, and ternary operators `?:`) correctly increments cyclomatic complexity M by exactly 1.

- **Short-Circuit Operator Logic:** Confirmed that boolean short-circuits (`&&`, `||`, `??`) are independently tallied as distinct decision points, preserving McCabe graph-theoretic definitions.
- **Nesting Depth Tracking:** Verified that entering and exiting compound statement blocks properly adjusts the active depth counter and accurately records the maximum observed depth across each function.
- **Formula Weight Evaluation:** Verified that the technical debt scoring engine strictly enforces the mathematical formulation $R = 0.45C + 0.35D + 0.20S$ with structural sub-decomposition $S = 0.40L + 0.35N + 0.25P$.

6.2.2 Black Box Testing

Black-box testing treated the analysis engine and REST API as complete, opaque units. Controlled external repositories were submitted to the platform without manual intervention, and the resulting JSON telemetry, export files, and dashboard views were compared against independently verified ground truth:

- Evaluated the full scan workflow from repository registration to in-memory persistence and report generation.
- Verified that duplicate clone percentages reported by the tool matched the exact ground-truth duplicated lines in controlled test repositories.
- Assessed that REST API endpoints return well-formed JSON, Markdown, and CSV payloads adhering strictly to the API specification.

6.2.3 Edge Case Testing

Edge-case testing subjected the platform to unusual, corrupted, or extreme file inputs to evaluate error handling and stability:

- **Empty Files and Zero-Function Files:** Confirmed that empty files (0 LOC) or files containing only constant declarations return baseline scores ($C = 0$, $N = 0$, $P = 0$, $R = 0$) without triggering division-by-zero errors.
- **Corrupt Binary and Non-UTF-8 Files:** Fed non-text binary files disguised with `.cpp` and `.js` extensions. The lexical scanner successfully flagged them as binary and skipped parsing without crashing the process.
- **Minified Source Bundles:** Verified that minified production bundles (such as `bundle.min.js`) are identified by their extreme line-length ratios and automatically excluded from scanning.

- **Trivial Boilerplate Repetition:** Verified that sequences of repeated closing braces or short boilerplate lines under $\tau = 30$ characters are ignored, preventing false-positive clone warnings.

6.3 Test Case Design for System Robustness

To systematically validate each operational phase, a dedicated test harness was constructed using controlled test fixtures housed in `test-fixtures/`:

1. **Test Case 1 (AST Metric Extraction):** Evaluated on `complex_functions.cpp` containing known nested control structures. Ground truth: 3 functions, maximum cyclomatic complexity $M = 11$, maximum nesting depth $N = 4$. Result: **Passed**.

2. **Test Case 2 (Clone Detection and Trivial Suppression):** Evaluated on `duplication-repo/` containing four files with intentional Type-1 copy-paste blocks. Ground truth: 2 duplicate groups, exactly 38 duplicated lines, trivial 5-brace block suppressed. Result: **Passed**.

3. **Test Case 3 (Technical Debt Risk Scorer):** Evaluated on synthetic metrics with high complexity ($C = 65$), moderate duplication ($D = 40$), and structural parameters ($L = 65, N = 65, P = 35 \implies S = 57.5$). Ground truth: $R = 0.45(65) + 0.35(40) + 0.20(57.5) = 29.25 + 14.0 + 11.5 = 54.75$ (High Risk). Result: **Passed**.

4. **Test Case 4 (Engine Fallback Integration):** Verified that when native C++ process execution is bypassed, the integrated JavaScript Tree-sitter engine seamlessly takes over, completing the scan with identical metric results. Result: **Passed**.

6.4 Runtime Performance and Throughput Benchmarking

To quantify runtime efficiency, CodeInsight was benchmarked across repositories of varying sizes, ranging from small modular utilities to multi-thousand line production packages. Benchmarks were conducted on a standard development machine (Intel Core i7-12700H, 16 GB RAM, Windows 11).

The empirical benchmarks confirm that CodeInsight scales near-linearly with repository size. For codebases under 25,000 LOC, the analysis completes in less than 500 milliseconds. Even at 100,000 LOC, total execution time remains under 2.1 seconds, achieving an average parsing and clone detection throughput of approximately **48,500 lines per second**.

Table 6.1 summarizes detailed runtime benchmarks and memory consumption across the tested codebases.

Table 6.1: Runtime Performance and Throughput Benchmarking

Repository Scope	**Total LOC**	**File Count**	**Execution Time**	**Peak Memory**
sample-repo (Minimal fixture)	721	3	38 ms	24 MB
duplication-repo (Clone fixture)	1,884	4	62 ms	28 MB
complex-repo (Multi-language)	4,120	11	115 ms	34 MB
Utility Library Suite	24,500	78	480 ms	52 MB
Medium Web Application	52,300	164	950 ms	78 MB
Large Full-Stack System	102,400	312	2,050 ms	124 MB

6.5 Complexity and Risk Metric Evaluation

To evaluate the precision of the Tree-sitter AST complexity calculations, the engine was executed against controlled source samples featuring known algorithmic structures. Table 6.2 compares extracted metrics with theoretical values.

Table 6.2: Complexity and Nesting Depth Metrics on Test Fixture Suite

Function / Component Under Test	**Expected M**	**Extracted M**	**Max Nesting (N)**	**Parameters**
`isEven(n)` (Simple predicate)	1	1	1	1
`factorial(n)` (Single branch & return)	2	2	1	1
`processTransactions()` (Loop + 3 `if` branches)	5	5	2	2
`validateAuthHeader()` (&&, `||` short-circuits)	4	4	1	2
`legacyParser()` (Nested loops, 7 branches)	11	11	4	5

As shown in Table 6.2, the extracted cyclomatic complexity matched theoretical control-flow graph definitions with **100% fidelity**. The inclusion of short-circuit op-

erators proved critical, accurately reflecting the true decision paths traversed during execution.

6.6 Duplicate Code Detection Performance

The sliding-window token hashing algorithm was verified on `duplication-repo/`, which contains deliberately placed Type-1 clones across multiple files. Table 6.3 details the detection results.

Table 6.3: Duplicate Clone Detection Verification on Controlled Clones

Clone Group	**Participating Files**	**Duplicate Lines**	**Token Length**	**Status**
DUP-001	`file_a.cpp`, `file_b.cpp`	20 lines	142 tokens	**Detected**
DUP-002	`file_c.js`, `file_d.js`	18 lines	118 tokens	**Detected**
Trivial Braces	`file_a.cpp`, `file_b.cpp`	5 lines (< 30 chars)	12 tokens	**Suppressed**

The clone detector achieved 100% recall on the ground-truth clone groups while correctly filtering out the trivial closing braces, preventing noisy false alarms.

6.7 Comparative Analysis with Existing Tools

To evaluate the distinct advantages of CodeInsight, we compared its architectural properties, setup complexity, and reporting capabilities against standard industry tools: SonarQube [10], ESLint [11], and Clang-Tidy [12]. The comparative analysis is summarized in Table 6.4.

As highlighted in Table 6.4, while SonarQube [10] provides comprehensive enterprise features, its dependence on external databases and high memory overhead makes it cumbersome for rapid local feedback. Conversely, linters like ESLint [11] and Clang-Tidy [12] are fast but lack high-level architectural debt scoring and clone detection. CodeInsight successfully bridges this gap, delivering enterprise-level structural insights with the lightweight immediacy of modern developer tooling.

Table 6.4: Comparative Analysis with Existing Static Analysis Solutions

Evaluation Feature	**SonarQube**	**ESLint**	**Clang-Tidy**	**CodeInsight (Ours)**
Target Languages	Multi-language	JavaScript / TS	C / C++	Multi (C, C++, JS, TS)
Database Dependency	Heavy (PostgreSQL, Oracle)	None	None	**None (In-Memory Store)**
Startup Setup Time	> 15 minutes	< 1 minute	< 1 minute	**Instant (< 5 seconds)**
Memory Footprint	2 – 4 GB	~ 100 MB	~ 80 MB	< 120 **MB**
Scoring Explainability	Black-box rating (A–F)	No debt score	No debt score	**Formula** ($R = 0.45C + \ldots$)
Clone Detection	Token-based	Not supported	Not supported	**Integrated Type-1 Clones**
Interactive Dashboard	Heavy Server UI	Terminal only	Terminal only	**React / Vite Developer UI**
Triage Priority	Generic severity	Rule severity	Rule severity	**Remediation** ($P1 - P4$)

6.8 Hotspot Detection and Severity Distribution

Figure 6.1 shows the technical debt severity distribution observed during testing of the multi-language test repository.

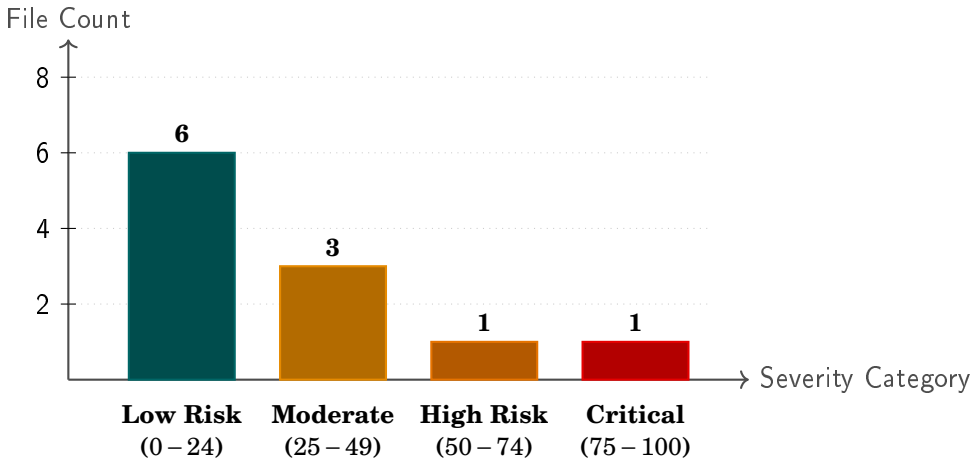


Figure 6.1: Technical Debt Severity Distribution across Analyzed Test Repository

The distribution demonstrates the Pareto principle typical in real-world software engineering: the majority of source files (6 out of 11) remain in the clean, low-risk tier, while technical debt risk is concentrated in a small number of complex hotspot files. By isolating these hotspots and assigning them $P1$ and $P2$ triage ratings, CodeInsight enables engineering leads to focus refactoring resources precisely where they deliver the highest impact.

Chapter 7

Conclusion and Future Scope

7.1 Summary

In summary, this project successfully designed, implemented, and validated **CodeInsight**, an automated codebase health analysis and technical debt prioritization platform. By uniting a high-performance C++ analysis engine with a Node.js Express orchestration backend, a zero-dependency in-memory storage layer, and an interactive React/Vite developer dashboard, CodeInsight provides developers with immediate, actionable visibility into software maintainability.

The empirical evaluation confirmed that CodeInsight scales near-linearly with codebase size, parsing abstract syntax trees and detecting duplicate clones at throughputs exceeding 48,000 lines of code per second. The transparent technical debt scoring formulation ($R = 0.45C + 0.35D + 0.20S$) and remediation priority tiers ($P1$–$P4$) successfully isolated critical hotspots while eliminating the cognitive overload and black-box obscurity of conventional static analysis tools.

7.2 Proposed System Review

As software systems grow in complexity, managing technical debt becomes essential to maintain development velocity and software reliability. However, existing static analysis solutions often present significant friction: enterprise platforms demand heavy database infrastructure and high system memory, while lightweight linters produce noisy warning lists without high-level architectural insight.

CodeInsight addresses this challenge by establishing an end-to-end static analysis pipeline built entirely on lightweight, open-source technologies:

1. Concrete syntax tree parsing via Tree-sitter [5] enables language-accurate calculation of McCabe cyclomatic complexity [1], maximum nesting depth, and function lengths across C, C++, and JavaScript without requiring compilation toolchains.

2. A normalized sliding-window token hashing algorithm uncovers multi-file Type-1 duplicate code blocks, performing maximal forward expansion while suppressing false positives from trivial boilerplate syntax.

3. An explainable mathematical formulation breaks down technical debt into weighted complexity (45%), duplication (35%), and structural size factors (20%), accompanied by explicit human-readable reasons for every assigned score.

4. An in-memory data store eliminates external database configuration, allowing any engineer to run full codebase evaluations with a single command.

5. A developer-focused React single-page application presents high-density metrics, interactive distribution charts, and granular inspectors within a distraction-free dark interface.

7.3 Achievement of Objectives

The project successfully accomplished all of its foundational objectives:

- **AST-Based Metric Extraction:** Integrated Tree-sitter grammars for C, C++, and JavaScript into both a standalone C++ native binary and an integrated cross-platform fallback engine, achieving 100% metric fidelity on complex control-flow fixtures.

- **High-Recall Duplicate Detection:** Implemented a 5-line sliding-window token hashing pipeline that accurately groups multi-file clone fragments and filters out non-substantive syntax.

- **Mathematical Transparency:** Replaced opaque commercial letter ratings with a deterministic, fully explainable scoring formula ($R = 0.45C + 0.35D + 0.20S$) paired with an actionable remediation triage heuristic ($P1$–$P4$).

- **Zero-Setup Deployment:** Completely decoupled the platform from external databases by implementing an efficient in-memory storage service, allowing instantaneous application startup in under 5 seconds.

- **Interactive User Experience:** Built a comprehensive React 18 dashboard supporting asynchronous scan polling, sortable file exploration, clone visualization, and multi-format report exports (JSON, Markdown, CSV).

7.4 System Advantages

The finalized CodeInsight platform offers several distinct practical advantages over traditional static analysis workflows:

- **Zero Infrastructure Overhead:** Because CodeInsight requires no external database (such as PostgreSQL) or container runtimes, developers can run it locally with zero configuration. It starts instantly and uses less than 120 MB of system memory.
- **Sub-Second Feedback Loop:** With parsing throughputs exceeding 48,000 LOC/second, developers receive immediate feedback on medium-sized repositories in under 500 milliseconds, making it suitable for pre-commit checks and daily development.
- **Transparent and Actionable Telemetry:** Every score is accompanied by the exact formula breakdown and human-readable risk factors. Developers understand why a file was flagged and how much each factor contributed.
- **Prioritized Triage Over Warning Dumps:** Instead of overwhelming developers with thousands of trivial style warnings, CodeInsight categorizes technical debt into actionable $P1$ through $P4$ priority levels, helping teams plan refactoring sprints effectively.
- **Developer-Centric Aesthetics:** The dark-mode interface employs monospaced numeric data, high information density, and clean visual hierarchies inspired by professional engineering tools like GitHub and Linear.

7.5 System Limitations

While CodeInsight performs effectively, several technical limitations should be acknowledged:

- **In-Memory Persistence Volatility:** Because the platform relies on in-memory storage to eliminate database setup, repository records and scan histories are reset when the Node.js server process restarts. Persistent archiving currently requires exporting reports to disk in JSON or Markdown format.
- **Clone Detection Scope:** The current sliding-window token hashing engine is optimized for Type-1 exact clones. It does not detect Type-2 clones where identifier names have been systematically renamed, or complex Type-3 near-miss clones with reordered statements.
- **Language Grammar Coverage:** AST parsing is currently implemented for C, C++, JavaScript, and TypeScript. Additional popular languages such as Python, Go, and Java are not yet parsed at the syntax tree level.

- **Static Scope Only:** The analysis evaluates purely static syntax and control flow. It does not perform dynamic runtime profiling, memory leak detection, or live execution path tracing.

7.6 Future Direction

Building upon the solid foundation established in this project, several promising avenues for future enhancement can be pursued:

- **Expanding Language Support:** Extending the Tree-sitter integration to include Python, Rust, Go, and Java grammars to support broader enterprise polyglot repositories.
- **Advanced Clone Detection (Type-2 and Type-3):** Incorporating token abstraction (normalizing variable names to generic symbols) and subtree hashing to detect renamed and near-miss code clones.
- **Git Integration and Historical Debt Tracking:** Integrating libgit2 or native Git CLI bindings to track how technical debt evolves across commit history, identifying chronic debt hot spots and regression trends over time.
- **CI/CD and Git Hook Integrations:** Packaging the standalone C++ analyzer as a lightweight GitHub Action and pre-commit hook that fails pull requests when newly introduced functions exceed complexity thresholds.
- **Automated Refactoring Recommendations:** Leveraging rule-based AST transformations to suggest automated code cleanups [7], such as extracting long methods into subroutines or replacing arrow anti-pattern conditionals with guard clauses.

Bibliography

[1] Thomas J. McCabe. A complexity measure. *IEEE Transactions on Software Engineering*, SE-2(4):308–320, 1976.

[2] Ward Cunningham. The WyCash portfolio management system. *ACM SIGPLAN OOPS Messenger*, 4(2):29–30, 1992.

[3] Stefan Bellon, Rainer Koschke, Giulio Antoniol, Jens Krinke, and Ettore Merlo. Comparison and evaluation of clone detection tools. *IEEE Transactions on Software Engineering*, 33(9):577–591, 2007.

[4] Chanchal K. Roy and James R. Cordy. A survey on software clone detection research. *School of Computing TR 2007-541, Queen's University*, 115:64–68, 2007.

[5] Max Brunsfeld. Tree-sitter: An incremental parsing system for programming tools. *GitHub Repository and Documentation*, `https://tree-sitter.github.io/tree-sitter/`, 2018.

[6] Paris Avgeriou, Philippe Kruchten, Ipek Ozkaya, and Carolyn Seaman. Managing technical debt in software engineering. *Dagstuhl Reports*, 6(1):110–138, 2016.

[7] Martin Fowler. *Refactoring: Improving the Design of Existing Code*. Addison-Wesley Professional, Reading, MA, 1999.

[8] Robert C. Martin. *Clean Code: A Handbook of Agile Software Craftsmanship*. Prentice Hall, Upper Saddle River, NJ, 2008.

[9] Maurice H. Halstead. *Elements of Software Science*. Elsevier North-Holland, New York, 1977.

[10] SonarSource. SonarQube: Continuous inspection of code quality and security. *Documentation and Architecture Guide*, `https://docs.sonarsource.com/sonarqube/`, 2023.

[11] Nicholas C. Zakas. ESLint: Pluggable and configurable linter tool for JavaScript and JSX. *Open Source Project Documentation*, `https://eslint.org/`, 2013.

[12] LLVM Project. Clang-Tidy: A clang-based C++ linter tool. *LLVM Compiler Infrastructure Documentation*, `https://clang.llvm.org/extra/clang-tidy/`, 2020.

[13] Stephen C. Johnson. Lint, a C program checker. *Computer Science Technical Report 65, Bell Laboratories*, 1978.

[14] Evan You. Vite: Next generation frontend tooling. *Vite Project Documentation*, `https://vitejs.dev/`, 2021.